VMLS Cheatsheet[1-9] - meanmachin3

1 Vectors

Vectors

•An ordered finite list of numbers.
•Block or stacked vectors($a = [b,c,d]$), Subvectors ($a_{r:s} = (a_r,...,a_s)$), Zero vectors (all elements equal to zero), Unit vectors(($e_i = 1$)), Ones vector(1_n) & Sparsity($nnz(x)$)

Vector addition

•*Commutative*: $a+b = b+a$
•*Associative*: $(a+b)+c = a+(b+c)$
•$a+0 = 0+a = a$
•$a-a = 0$

1.1 Scalar-vector multiplication

$(-2)(1,9,6) = (-2,-18,-12)$
•*Commutative*: $\alpha a = a\alpha$
•*Left-distributive*: $(\beta+\gamma)a = \beta a+\gamma a$
•*Right-distributive*: $a(\beta+\gamma) = a\beta+a\gamma$

Linear combinations: $\beta_1 a_1 + ... + \beta_m a_m$
•With Unit vectors: $b = b_1 e_1 + ... + b_n e_n$
•If $\beta_1 + ... + \beta_m = 1$, linear combination is said to be *affine combination*

1.2 Inner product

$a^T b = a_1 b_1 + a_2 b_2 + ... + a_n b_n$ **Properties**:
•*Commutativity*:$a^T b = b^T a$
•*Scalar multiplication Associativity*: $(\gamma a)^T b = \gamma(a^T b)$
•Vector addition Distributivity: $(a+b)^T c = a^T c + b^T c$.

General examples:

•Unit vector: $e_i^T a = a_i$
•Sum: $\mathbf{1}^T a = a^1 + ... + a^n$
•Average: $(\mathbf{1}/n)^T a = (a^1 + ... + a^n)/n$
•Sum of squares: $a^T a = a_1^2 + ... + a_n^2$
•Selective sum: If $b_i = 1 or 0$, $b^T a$ is the sum of elements for which $b_i = 1$,

Block vectors

$a^T b = a_1^T b_1 + ... + a_k^T b_k$

1.3 Complexity of vector computations

•*Space*: 8n bytes
•*Complexity of vector operations*: $x^T y = 2n-1$ flops (n scalar multiplications and $n-1$ scalar additions)
•*Complexity of sparse vector operations*: If x is sparse, then computing ax requires **nnz**(x) flops, If x and y are sparse, computing x + y requires no more than min**nnz**(x), **nnz**(y). computing $x_T y$ requires no more than 2 min**nnz**(x), **nnz**(y) flops

2 Linear functions

2.1 Linear functions

$f : R^n \to R$ means f is a function mapping n-vectors to numbers
Superposition & linearity: $f(\alpha x+\beta y) = \alpha f(x)+\beta f(y)$
•$f(\alpha_1 x_1 + ... + \alpha_k x_k) = \alpha_1 f(x_1) + ... + \alpha_k f(x_k)$
A function that satisfies superposition is called *linear*

Linear function satisfies

•*Homogeneity*: For any n-vector x and any scalar α, $f(\alpha x) = \alpha f(x)$
•*Additivity*: For any n-vectors x and y, $f(x+y) = f(x)+f(y)$

Affine functions $f : R_n \to R$ is affine if and only if it can be expressed as $f(x) = a^T x + b$ for some n-vector a and scalar b, which is sometimes called the *offset*
•Any *affine* scalar-valued function satisfies the following variation on the superposition property: $f(\alpha x+\beta y) = \alpha f(x)+\beta f(y)$, where $\alpha+\beta = 1$

2.2 Taylor approximation

The (first-order) Taylor approximation of f near (or at) the point z:

$\hat{f}(x) = f(z) + \frac{\partial f}{\partial x_1}(z)(x_1 - z_1) + ... + \frac{\partial f}{\partial x_n}(z)(x_n - z_n)$

Alternatively, $\hat{f}(x) = f(z) + \nabla f(z)^T (x-z)$

2.3 Regression model

Regression model is (the affine function of x) $\hat{y} = x^T\beta + v$

3 Norm and distance

3.1 Norm

Euclidean norm (or just norm) is

$\|x\| = \sqrt{x_1^2 + x_2^2 + ... + x_n^2} = \sqrt{x^T x}$

Properties

•homogeneity: $\|\beta x\| = |\beta|\|x\|$
•triangle inequality: $\|x+y\| \le \|x\| + \|y\|$
•non negativity: $\|x\| \ge 0$
•definiteness: $\|x\| = 0$ only if x = 0
positive definiteness = non negativity + definiteness

rms(x) = $\sqrt{\frac{x_1^2+...+x_n^2}{n}} = \frac{\|x\|}{\sqrt{n}}$

•*Norm of a sum*:
$\|a+b\|^2 = (x+y)^T(x+y) = \|x\|^2 + 2x^T y + \|b\|^2$
Norm of block vectors $\|(a,b,c)\| = \sqrt{\|a\|^2 + \|b\|^2 + \|c\|^2} = \|(\|a\|, \|b\|, \|c\|)\|$
Chebyshev inequality k of its entries satisfy $|x_i| \ge a$,
then $\frac{k}{n} \le (\frac{\mathbf{rms}(x)}{a})^2$

3.2 Distance

dist$(a,b) = \|a-b\|$
Triangle Inequality: $\|a-c\|^2 = \|(a-b)+(b-c)\| \le \|a-b\| + \|b-c\|$
z_j is the nearest neighbor of x if $\|x-z_j\| \le \|x-z_i\|, i = 1,..,m$

3.3 Standard Deviation

de-meaned vector: $\tilde{x} = x - \mathbf{avg}(x)\mathbf{1}$
standard deviation:

std(x) = **rms**$(\tilde{x}) = \frac{\|x-(\mathbf{1}^T x/n)\mathbf{1}\|}{\sqrt{n}}$

$\mathbf{rms}(x)^2 = \mathbf{avg}(x)^2 + \mathbf{std}(x)^2$

By Chebyshev inequality, $|x_i - \mathbf{avg}(x)| \ge \alpha\mathbf{std}(x)$ then $k/n \le (std(x)/a)^2$. (This inequality is only interesting for $a > std(x)$)
Cauchy–Schwarz inequality: $|a^T b| \le \|a\|\|b\|$

3.4 Angle

angle between two nonzero vectors a, b defined as
$\angle(a,b) = arccos(\frac{a^T b}{\|a\|\|b\|})$
$a^T b = \|a\|\|b\|cos(\angle(a,b))$

Classification of angles

$\theta = \pi/2$: $a \perp b$
$\theta = 0$: $a^T b = \|a\|\|b\|$
$\theta = \pi = 180° : a^T b = -\|a\|\|b\|$
$\theta \le \pi/2 = 90° = a^T b \ge 0$
$\theta \ge \pi/2 = 90° = a^T b \le 0$

Correlation Coeficient (ρ) $\rho = \frac{\tilde{a}^T\tilde{b}}{\|\tilde{a}\|\|\tilde{b}\|}$

With $u = \tilde{a}/\mathbf{std}(a)$ & $u = \tilde{b}/\mathbf{std}(b)$,
$\rho = u^T v/n$ where $\|u\| = \|v\| = n$

$std(a+b) = \sqrt{std(a)^2 + 2\rho std(a)std(b) + std(b)^2}$

Properties of standard deviation

•**std**$(x+a1) =$ **std**(x)
•**std**$(ax) = |a|std(x)$

Standardization $z = \frac{1}{\mathbf{std}(x)}(x - \mathbf{avg}(x)\mathbf{1})$

3.5 Complexity

•*norm*: 2n
•***rms***: 2n
•***dist***(a,b): 3n
•$\angle(a,b)$: 6n

4 Clustering

4.1 A clustering Objective

$G_j \subset \{i|c_i = j\}$ where G_j is set of all indices i for which $c_i = j$
•Group representatives: n-vectors $z_1,..,z_k$
•Clustering objective is
$J^{clust} = \frac{1}{N}\sum_{i=1}^N \|x_i - Z_{c_i}\|^2$
•mean square distance from vectors to associated representative
•goal: choose clustering c_i and representatives z_j to minimize J_clust

4.2 The k-means algorithm

---

given $x_1,..,x_N \in \mathbf{R}^n$ and $z_1,...,z_k \in \mathbf{R}^n$
repeat
– *Update partition*: assign i to G_j, $j = argmin_{j'}\|x_i - z_{j'}\|_2$
– *Update centroids*: $Z_j = \frac{1}{|G_j|}\sum_{i\in G_j} x_i$
until $z1,...,zk$ stop changing

---

5 Linear Independence

$(a_1,...,a_k)$ is linearly dependent if $\beta_1 a_1 + ... + \beta_k a_k = 0$, for some $\beta_1,,...\beta_k$,that are not all zero

5.1 Linear Independence

$(a_1,...,a_k)$ is linearly independent if $\beta_1 a_1 + ... + \beta_k a_k = 0$ & $\beta_1 = ... = \beta_k = 0$
•Adding vector to linearly dependent makes new vector linearly dependent
•Removing vector from linearly independent makes new vector linearly independent

5.2 Basis

basis: A collection of n linearly independent(maximum possible size) n-vectors

Independence-dimension inequality

•*a linearly independent set of n-vectors can have at most n elements*
•*any set of n + 1 or more n-vectors is linearly dependent*

5.3 Orthonomal Vectors

$a_1,...,a_k$ are (mutually) *orthogonal* if $a_i \perp a_j$ for i != j
They are *normalized* if $\|a_i\| = 1$ for i=1,..,k
•*orthonormal* if *orthogonal* & *normalized*
•can be expressed using inner products
$a_i^T a_j = \begin{cases}1, & \text{if } i = j\\ 0, & i \ne j\end{cases}$
•orthonormal sets of vectors are linearly independent
•$a_1,...,a_n$ is an orthonormal basis, we have for any n-vector $x = (a_1^T x)a_1 + ... + (a_n^T x)a_n$

5.4 Gram–Schmidt(orthogonalization)

An algorithm to check if $a_1,...,a_k$ are linearly independent

---

given n-vectors $a_1,...a_n$
for i = 1,..,k
1.Orthogonalization:
$\tilde{q}_i = a_i - (q_1^T a_i)q_1 - ... - (q_{i-1}^T a_i)q_{i-1}$
2. Test for linear dependence:
if $\tilde{q} = 0$, quit
3.Normalization: $q_i = \tilde{q}_i/\|\tilde{q}_i\|$

•if G–S does not stop early (in step 2), $a_1,...,a_k$ are linearly independent
•if G–S stops early in iteration $i = j$, then a_j is a linear combination of $a_1,..,a_{j-1}$ (so $a_1,..,a_k$ are linearly dependent)

Complexity: $2nk^2$

6 Matrices

6.1 Matrices

The set of real m × n matrices is denoted $R^{m\times n}$

6.2 Zero and identity matrices

•*Zero*: All elements equals 0.
•*Identity*: All elements equals 0 and diagonal element equals 1.
•*Sparse*: If many entries are 0
•*Diagonal*: off-diagonal entries are zero
•*Triangular*: *upper triangular* if $A_{ij} = 0$ for $i > j$, and it is *lower triangular* if $A_{ij} = 0$ for $i < j$
•*Adjacency Matrix*:
For, $R = (1,2),(1,3),(2,1),(2,4),(3,4),(4,1)$
$A_{ij} = \begin{cases}1, & (i,j) \in R\\ 0, & (i,j) \notin R\end{cases}$
A relation R on $1,...,n$ is represented by the n×n matrix A with $A_{ij} = 1$, if there exists a edge else , $A_{ij} = 0$

6.3 Transpose, addition and norm

Block matrix Transpose

$\begin{bmatrix}A & B\\ C & D\end{bmatrix}^T = \begin{bmatrix}A^T & C^T\\ B^T & D^T\end{bmatrix}$

Symmetric matrix: $A = A^T$

Properties of matrix addition

•Commutativity: $A+B = B+A$
•Associativity: $(A+B)+C = A+(B+C)$
•Addition with zero matrix: $A+0 = 0+A = A$
•Transpose of sum: $(A+B)^T = A^T + B^T$
If A is a matrix and β , γ are scalars $(\beta+\gamma)A = \beta A + \gamma A, (\beta\gamma)A = \beta(\gamma A)$

Matrix norm $\|A\| = \sqrt{\sum_{i=1}^n\sum_{j=1}^m A_{ij}^2}$ matrix norm satisfies the properties of any norm

6.4 Matrix-vector multiplication

A is an $m \times n$ matrix and x is an n-vector, then the matrix-vector product $y = Ax$

$y_i = \sum_{k=1}^n A_{ik}x_k = A_{i1}x_1 + ... + A_{in}x_n$ for $i = 1...m$

•Row and column interpretations.

$y = Ax$ can be expressed as $y_i = b_i^T x, i = 1,..,m$ where $b_1^T,...,b_m^T$ are rows of A
•$y = Ax$ could also be expressd in terms of column $y = x_1a_1 + x_2a_2 + ... + x_na_n$

General Examples

•Picking out columns and rows An important identity is $Ae_j = a_j$, the jth column of A. (In other words, $(A^T e_i)^T$ is the ith row of A.)
•*Summing or averaging columns or rows*: The m-vector A1 is the sum of the columns of A; its ith entry is the sum of the entries in the ith row of A. The m-vector $A(\mathbf{1}/n)$ is the average of the columns of A; its ith entry is the average of the entries in the ith row of A. In a similar way, $A^T\mathbf{1}$ is an n-vector, whose jth entry is the sum of the entries in the jth column of A.

6.5 Complexity

addition: mn
sparse matrix addition: If A or B or both are sparse $min\{nnz(A), nnz(B)\}$
vector multiplication A_{mxn} *with n-vector*: $m(2n-1) \approx 2mn$
Matrix Transpose: 0 flops

7 Matrix examples

7.1 Geometric transformations

•*Scaling*: $y = Ax$ with $A = aI$ stretches a vector by the factor $|a|$ (or shrinks it when $|a| < 1$), and it flips the vector (reverses its direction) if $a < 0$

•*Dilation*: $y = Dx$, where D is a diagonal matrix, $D = diag(d1, d2)$. Stretches the vector x by different factors along the two different axes. (Or shrinks, if $|d_i| < 1$, and flips, if $di < 0$.)

•*Rotation Matrix (counter clockwise)*: $y = \begin{bmatrix} cos\theta & -sin\theta \\ sin\theta & cos\theta \end{bmatrix} x$

•*Reflection* Suppose that y is the vector obtained by reflecting x through the line that passes through the origin, inclined θ radians with respect to horizontal.

$y = \begin{bmatrix} cos(2\theta) & sin(2\theta) \\ sin(2\theta) & -cos(2\theta) \end{bmatrix} x$

•*Projection into a line* Projection of point x onto a set is the point in the set that is closest to x.

$y = \begin{bmatrix} (1/2)(1+cos(2\theta)) & (1/2)sin(2\theta) \\ (1/2)sin(2\theta) & (1/2)(1-cos(2\theta)) \end{bmatrix} x$

7.2 Selectors

An $m \times n$ selector matrix A is one in which each row is a unit vector (transposed):

$$\begin{bmatrix} e_{k_1}^T \\ . \\ . \\ e_{k_m}^T \end{bmatrix}$$

When it multiplies a vector, it simply copies the k_ith entry of x into the ith entry of $y = Ax$:

$y = (x_{k_1}, x_{k_2}, ..., x_{k_m})$

r:s matrix slicing

$A = [0_{m \times (r-1)} I_{m \times m} 0_{m \times (n-s)}]$

where $m = s - r + 1$

7.3 Incidence matrix

Directed graph: A *directed graph* consists of a set of *vertices* (or nodes), labeled 1,...,n, and a set of *directed edges* (or branches), labeled 1,...,m.

$A_{ij} = \begin{cases} 1, & \text{edge j points to node i} \\ -1, & \text{edge j points from node i} \\ 0, & \text{otherwise} \end{cases}$

7.4 Convolution

The convolution of an n-vector a and an m-vector b is the (n + m - 1)-vector denoted c = a * b

$c_k = \sum_{i+j=k+1} a_i b_j, k = 1, .., n+m-1$

Properties of convolution

•symmetric: $a * b = b * a$

•associative: $(a * b) * c = a * (b * c)$

•$a * b = 0$ implies that either $a = 0$ or $b = 0$

•A basic property is that for fixed a, the convolution a * b is a linear function of b; and for fixed b, it is a linear function of a, $a * b = T(b)a = T(a)b$ where where T(b) is the (n + m - 1) × n matrix with entries

$T(b)_{ij} = \begin{cases} b_{i-j+1}, & 1 \le i-j+1 \le m \\ 0, & \text{otherwise} \end{cases}$

Complexity of convolution

•$c = a * b$: 2mn flops

•$T(a)bor T(b)a$: 2mn flops

•Convolution could be calculated faster using *fast Fourier transform (FFT)* : $5(m+n)log_2(m+n)flops$

8 Linear equations

8.1 Linear and affine functions

•Superposition condition:$f(\alpha x + \beta y) = \alpha f(x) + \beta f(y)$

•Such an f is called Linear

Matrix vector product function:

•A is $m \times n$ matrix such that $f(x) = Ax$

•f is linear: $f(\alpha x + \beta y) = A(\alpha x + \beta y) = \alpha f(x) + \beta f(y)$

•Converse is true: If $f : R^n \mapsto R^m$ is linear, then

$f(x) = f(x_1 e_1 + x_2 e_2 + ...x_n e_n)$

$= x_1 f(e_1) + x_2 f(e_2) + ...x_n f(e_n) = Ax$ with $A = [f(e_1) + f(e_2) + ...f(e_n)]$

Affine Functions: $f : R^n \mapsto R^m$ is affine if it is a linear function plus a constant i.e $f(x) = Ax + b$ same as $f(\alpha x + \beta y) = \alpha f(x) + \beta f(y)$ holds for all x, y and α, β such that $\alpha + \beta = 1$

A and b can be calculated as

$A = [f(e_1) - f(0)\ \ f(e_2) - f(0)...f(e_n) - f(0)]]$;

$b = f(0)$

•Affine functions sometimes incorrectly called linear functions

8.2 Linear function models

Price elasticity of demand $\delta_i^{price} = (p_i^{new} - p_i)/p_i$: fractional changes in prices

$\delta_i^{dem} = (d_i^{new} - d_i)/d_i$: fractional change in demand Price demand elasticity model: $\delta^{dem} = E\delta^{price}$

Taylor series approximation

•The (first-order) Taylor approximation of f near (or at) the point z:

$\hat{f}(x) = f(z) + \frac{\partial f}{\partial x_1}(z)(x_1 - z_1) + ... + \frac{\partial f}{\partial x_n}(z)(x_n - z_n)$

•in compact notation:

$\hat{f}(x) = f(z) + Df(z)(x - z)$

8.3 Systems of linear equations

•set (or system) of m linear equations in n variables $x_1, .., x_n$:

$A_{11}x_1 + A_{12}x_2 + ... + A_{1n}x_n = b_1$

$A_{21}x_1 + A_{22}x_2 + ... + A_{2n}x_n = b_2$

.

.

$A_{m1}x_1 + A_{m2}x_2 + ... + A_{mn}x_n = b_m$

•systems of linear equations classified as

– under-determined if m < n (A wide)

– square if m = n (A square)

– over-determined if m > n (A tall)

Balancing equation example

•consider reaction with m types of atoms, p reactants, q products

•$m \times p$ reactant matrix R is defined by R_{ij} = number of atoms of type i in reactant R_j

for i = 1,..,m and j = 1,..,p

•with $a = (a_1, .., a_p)$ (vector of reactant coefficients)

Ra = (vector of) total numbers of atoms of each type in reactants

•define product $m \times q$ matrix P in similar way

•m-vector Pb is total numbers of atoms of each type in products

•conservation of mass is $Ra = Pb$

•conservation of mass is $[R - P][a\ b]^T = 0$

•simple solution is $a = b = 0$

•to find a nonzero solution, set any coefficient (say, a_1) to be 1

•balancing chemical equations can be expressed as solving a set of m + 1 linear equations in p + q variables

$\begin{bmatrix} R & -P \\ e_1^T & 0 \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = e_{m+1}$

(we ignore here that a_i and b_i should be nonnegative integers)

9 Linear dynamical systems

9.1 Linear dynamical systems

$x_{t+1} = A_t x_t$, t = 1,2,. . .

•A_t are n × n dynamics matrices

•$(A_t)_{ij}(x_t)_j$ is contribution to $(x_{t+1})_i$ from $(x_t)_j$

•system is called time-invariant if $A_t = A$ doesn't depend on time

•can simulate evolution of xt using recursion $x_{t+1} = A_t x$

•linear dynamical system with input $x_{t+1} = A_t x_t + B_t u_t + c_t$, t = 1,2,...

– u_t is an input m-vector

– B_t is n × m input matrix

– c_t is offset

K-Markov model:

$x_{t+1} = A_1 x_t + ... + A_K x_{t-K+1}$, t = K,K + 1,...

– next state depends on current state and K - 1 previous states

– also known as auto-regresssive model

– for K = 1, this is the standard linear dynamical system $x_{t+1} = Ax_t$

9.2 Population dynamics

•$x_t \in R^{100}$ gives population distribution in year t = 1,...,T

•$(x_t)^i$ is the number of people with age i - 1 in year t (say, on January 1)

•birth rate $b \in R^{100}$, death (or mortality) rate $d \in R^{100}$

•b_i is the number of births per person with age i - 1

•d_i is the portion of those aged i - 1 who will die this year (we'll take $d_{100} = 1$)

•let's find next year's population distribution $x_t + 1$ (ignoring immigration)

•number of 0-year-olds next year is total births this year:

$(x_{t+1})_1 = b^T x_t$

•number of i-year-olds next year is number of (i - 1)-year-olds this year, minus those who die: $(x_{t+1})_{i+1} = (1 - d_i)(x_t)_i$, i = 1,..,99

•$x_{t+1} = Ax_t$, where

$A = \begin{bmatrix} b_1 & b_2 & ... & b_{99} & b_{100} \\ 1-d_1 & 0 & ... & 0 & 0 \\ 0 & 1-d_2 & ... & 0 & 0 \\ . & . & ... & . & . \\ . & . & ... & . & . \\ 0 & 0 & ... & 1-d_{99} & 0 \end{bmatrix}$

9.3 Epidemic dynamics

SIR Model

•4-vector x_t gives proportion of population in 4 infection states

– *Susceptible*: can acquire the disease the next day

– *Infected*: have the disease – *Recovered*: had the disease, recovered, now immune

– *Deceased*: had the disease, and unfortunately died

•sometimes called SIR model

•e.g., $x_t = (0.75, 0.10, 0.10, 0.05)$ **over each day**,

•among susceptible population,

– 5% acquires the disease

– 95% remain susceptible •among infected population,

– 1% dies

– 10% recovers with immunity

– 4% recover without immunity (i.e., become susceptible)

– 85% remain infected

•100% of immune and dead people remain in their state

•epidemic dynamics as linear dynamical system

$x_{t+1} = \begin{bmatrix} 0.95 & 0.04 & 0 & 0 \\ 0.05 & 0.85 & 0 & 0 \\ 0 & 0.10 & 1 & 0 \\ . & . & ... & . \\ . & . & ... & . \\ 0 & 0.01 & 0 & 1 \end{bmatrix} x_t$